PHILOWAY CONSULTING CHALLENGE

Adopt AI to release audit capacity while protecting client trust

Cloud, on-premise, hybrid or wait? A recommendation with its business case, its risk boundary and its implementation plan.

CONSULTING TEAM

Presenter A Presenter B Presenter D Presenter C Silas Zamzow

GROUP

Cohort B, Group 4

EDHEC BUSINESS SCHOOL

Q1 STRATEGIC ASSESSMENT

Why invest in AI and why do it now?

WHY INVEST IN AI, AND WHAT BUSINESS VALUE?

RELEASE CAPACITY

Search and preparation absorb hours clients do not pay for

DEFEND THE SERVICE

Faster, evidence-backed findings at unchanged independence

OWN THE KNOWLEDGE

Our case knowledge is what rivals cannot buy

COST OF WAITING

The prerequisites take calendar time whenever we decide to start

WHAT IS THE RISK OF NOT INVESTING NOW?

- Peers already run internal assistants at far larger scale
- Every client must sign an ICT agreement before use
- Our accreditation body expects notice before first deployment
- Auditors are now expected to understand AI-using clients

Waiting does not postpone the prerequisites. It only delays the pay-off.

Q1 STRATEGIC ASSESSMENT

Our peers already run their own AI assistants, and they built them the way we propose

1

HOW ARE COMPETITORS ALREADY INTEGRATING AI?

- Two sector peers run their own internal assistants
- Retrieval over standards, past reports and prior findings
- Used for preparation, drafting findings and knowledge reuse
- Deployed on business cloud rather than local hardware
- Internal use only, under documented confidentiality controls

WHAT DOES THIS IMPLY FOR A TWENTY-PERSON FIRM?

- No competitor published a measured productivity return
- The architecture is proven, the business case is ours
- We compete on internal efficiency
- Auditor AI competence sits in our scope and costs training
- Becoming an accredited ISO 42001 body is out of reach

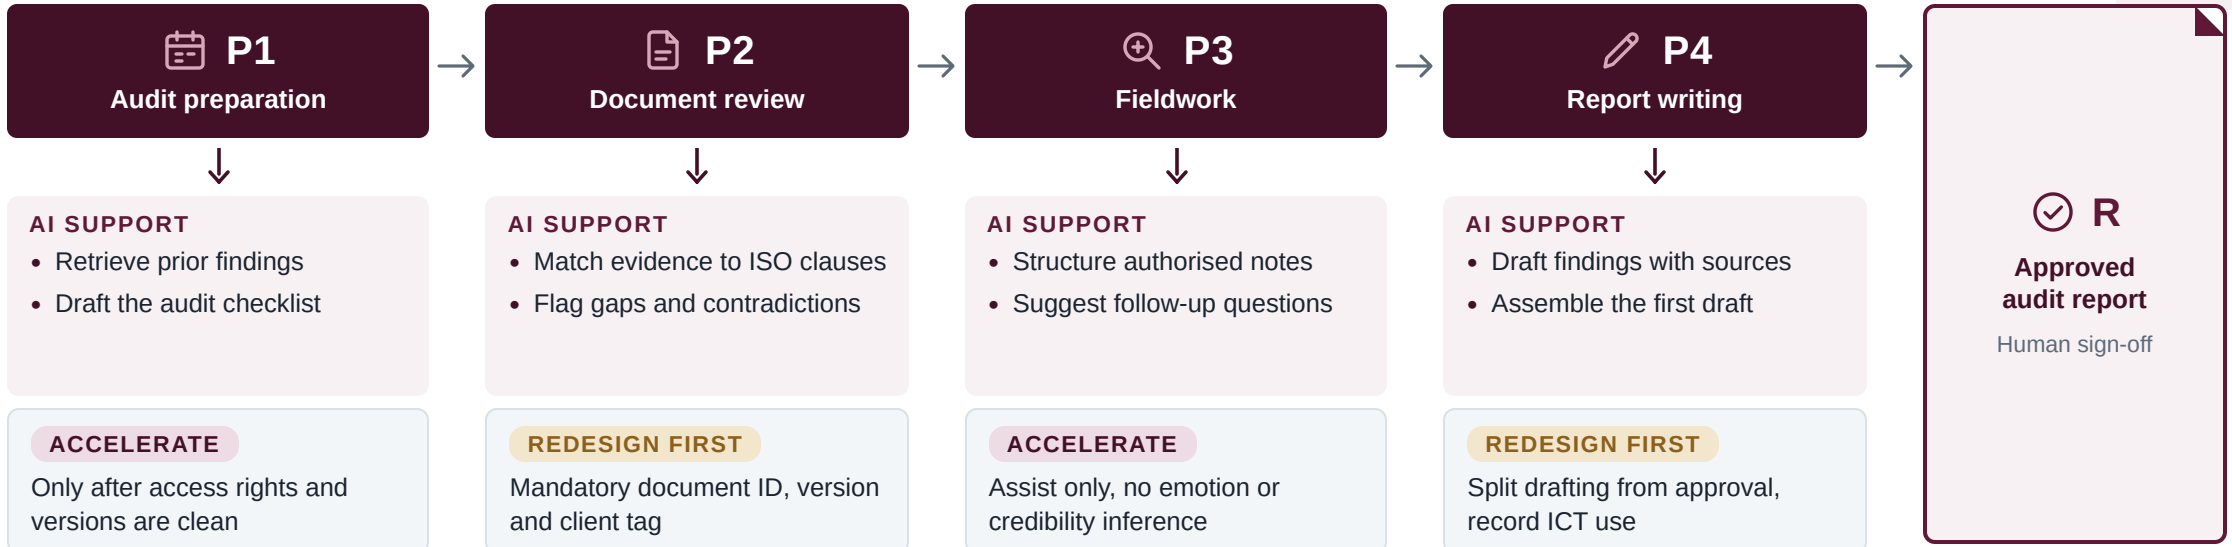
Not the same thing: SGS, DEKRA, TÜV, DNV and LRQA now sell certification *of* client AI systems. Selling assurance about AI is a service decision. Running an assistant inside the audit process is the operational benchmark, and that is the one we are measuring ourselves against.

Nobody is ahead on measured productivity. The race is on lead time.

Q2 OPERATIONS AND PROCESS IMPACT

AI supports each audit step, while auditors remain accountable for the final report

HOW COULD AI AFFECT EACH STEP, AND MUST THE STEP BE REDESIGNED FIRST?



WHAT HAS TO EXIST BEFORE ANY OF THIS WORKS?

- Knowledge management has to come first
- A curated, permissioned base must exist before retrieval works
- The audit RAG searches that base and cites every source
- Internal support starts the pilot because its content is approved

Q3 CONFIDENTIALITY AND RISK ANALYSIS

Risk governs which documents enter the system

WHICH RISKS ARISE, AND HOW MUCH SHOULD EACH ONE WEIGH?

	RISK	WEIGHT	WHAT IT ACTUALLY DECIDES
	Confidentiality and reputation	■■■	Which documents may leave the building. One serious breach ends the client relationship.
	Regulatory and contractual	■■■	Written client agreement before use, manual fallback on refusal, DPIA before indexing.
	Incorrect findings	■■■	Mandatory human approval and refusal without evidence. A citation is not a conclusion.
	Cybersecurity	■■□	Standard controls either way. Local adds patching and insider duties rather than removing them.
	Vendor dependency	■□□	Exportable documents and metadata, tested switching, annual contract review.

HOW MUCH SHOULD RISK WEIGH AGAINST THE VALUE AT STAKE?

- Fines are capped; accreditation loss is uncapped
- Every excluded document costs us benefit
- Local removes the provider, adds insider access
- A weighted scorecard would fake precision
- Exclude first, then compare the compliant options

Not one of these five risks selects cloud over local. They select **what enters the system and who signs it off.**

Q4 TECHNOLOGY OPTIONS

Cloud wins five of six criteria; on-premise wins only where nothing may leave the building

HOW DO CLOUD AND ON-PREMISE COMPARE ON THE SAME CRITERIA?

CRITERION	CLOUD	ON-PREMISE	BETTER FIT
Speed to value	Pilot running in weeks on existing licences	Procurement, build and tuning before first use	Cloud
Scalability	Capacity bought in small increments, per user	Grows in hardware steps; one server is not high availability	Cloud
Data protection	Processing leaves the firm, bounded by contract and region	Inference and index stay on the internal network	On-premise
User experience	Familiar interface, maintained and updated by the provider	Own portal to build, own quality and latency to tune	Cloud
Maintenance	Provider patches models and platform; we own content	We own patching, backup, restore tests and capacity	Cloud
Cost at our scale	EUR 77,424 over three years	EUR 131,000 over three years	Cloud

The honest caveat: on-premise does not remove risk, it moves it. The provider disappears, patching, insider access, backups and outages arrive. In a twenty-person firm total risk can rise.

Hybrid is not a compromise: it needs both operating capabilities, adds routing errors and costs EUR 173,164. It becomes the answer only if a client contractually requires local processing.

Q4 TECHNOLOGY OPTIONS

Both paths are buildable today with named products, and only one is reversible

4

WHICH REAL SUPPLIERS, CONFIGURATIONS AND INDICATIVE COSTS?

CLOUD CONFIGURATION

- **Microsoft 365 Copilot Business** for approved office work
- **Azure OpenAI**, GPT-4.1-mini class, EU Data Zone
- **PostgreSQL with pgvector**, EU-hosted audit index
- **Entra ID portal**, permissions enforced before retrieval
- Three free settings: no web search, no flex routing, EU pinned

EUR 18.20/user/month **EUR 120**/month model **EUR 200**/month hosting

List prices and budget allowances, confirm before purchase

ON-PREMISE CONFIGURATION

- **Dell Precision 7960**, 128 GB RAM, 4 TB SSD
- **NVIDIA RTX PRO 6000**, 96 GB, plus UPS and backup
- **Qwen3-32B** under Apache 2.0, served with vLLM
- **PostgreSQL with pgvector** and local OCR
- VPN access and a contracted IT operator

EUR 25,000 hardware package **EUR 0** model licence, Apache 2.0

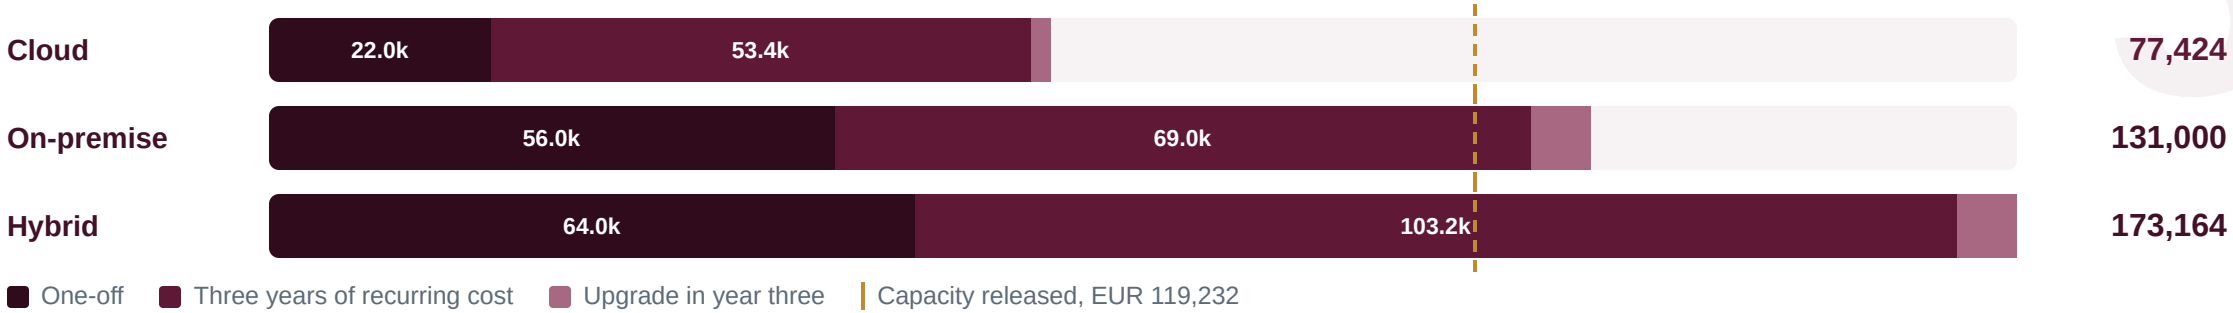
Hardware subject to quotation, no list price available

The cloud configuration can be switched off; the hardware cannot be unbought. We keep a **current quote on file** so the on-premise path stays selectable at month twelve. **What these components add up to is the next slide.**

Q5 FINANCIAL ANALYSIS

Only the cloud route stays below the value it is meant to create

WHAT DOES EACH OPTION COST OVER THREE YEARS?



WHERE DOES THE MONEY GO?

ONE-OFF INVESTMENT	CLOUD	ON-PREMISE	HYBRID
Hardware and infrastructure	0	25,000	25,000
Integration and RAG build	12,000	18,000	24,000
Data, security and contracts	6,000	8,000	9,000
Training and staff learning time	4,000	5,000	6,000
Total before go-live	**22,000**	**56,000**	**64,000**

EVERY YEAR	CLOUD	ON-PREMISE	HYBRID
Licence subscriptions	4,368	0	4,368
Model usage	1,440	0	720
Hosting, power and backup	2,400	1,800	2,400
External maintenance and support	3,600	7,200	8,400
Internal ownership	4,000	12,000	16,000
Governance and quality assurance	2,000	2,000	2,500
Total per year	**17,808**	**23,000**	**34,388**

Q6 ORGANIZATIONAL READINESS

Twenty people can own a managed platform; nobody here can run one

WHICH CAPABILITIES DOES EACH OPTION NEED, AND WHO WOULD OWN THEM?

CAPABILITY	TODAY	NEEDED FOR CLOUD	NEEDED FOR ON-PREMISE	OWNER
Audit domain knowledge	Have it	Unchanged. The scarce asset stays internal.	Unchanged	Audit and quality lead
Sources, tests and approval rules	Partly	0.05 FTE, about two hours per week	0.15 FTE, plus content duties	Internal product owner
Platform operation	None	Bought from the provider and an IT partner	Own patching, backup, restore tests, capacity	External IT partner
Privacy and accreditation	Partly	DPIA, processor terms, client ICT annex	The same, plus insider-access controls	Privacy function
Availability and cover	None	Provider service levels	One server is not high availability	Not achievable in-house

ARE THE PROCESSES READY TO ABSORB AI?

- Internal support and preparation can start as they are
- Document review and reporting need redesign first
- The knowledge base must be curated before retrieval

WHAT RESISTANCE SHOULD WE EXPECT, AND WHAT DO WE DO?

- Two practitioners co-design the test set from week one
- Released time goes to quality before it goes to targets
- Excess access rights are corrected before switch-on

Q7 FUTURE SCALABILITY

The recommendation survives growth; the audit workflow does not survive it unchanged

WOULD THE SOLUTION STILL WORK AT TRIPLE VOLUME, DOUBLE HEADCOUNT AND ABROAD?

Three times the documents	Twice the headcount	International expansion	The binding constraint
SOLUTION Incremental indexing, duplicate detection and version control.	**SOLUTION** Forty licences and a load test at ten concurrent requests.	**SOLUTION** Country, language and standard metadata, tested per language.	**SOLUTION** Technology scales in small increments.
ORGANISATION Source ownership assigned per standard and per client.	**ORGANISATION** Approval rights spread across teams instead of one owner.	**ORGANISATION** Local reviewers and local accreditation relationships.	**ORGANISATION** Expert review capacity must grow with throughput.
PROCESS Retention and deletion rules become mandatory.	**PROCESS** Onboarding carries AI training as a standard module.	**PROCESS** Client agreements renegotiated under each jurisdiction.	**PROCESS** Faster research without more reviewers only moves the queue.

At forty people and thirty thousand documents, cloud remains the right direction. The annual run rate rises to roughly **EUR 35,816**.

Q8 IMPLEMENTATION STRATEGY

A twelve-week pilot on approved content buys the evidence for every later decision

8

WHICH LOW-RISK PILOT, AND WHAT HAPPENS SHORT, MEDIUM AND LONG TERM?

WEEKS 1 TO 2

Baseline times measured, data classes sized in hours, ICT annex drafted and the accreditation body informed.

Gate Scope, named owners and a documented eligibility estimate agreed

WEEKS 3 TO 6

The pilot use case: internal procedures and synthetic cases. Five users and no real client file.

Gate Fifty validated questions, ten unanswerable, access tests passed

WEEKS 7 TO 12

After written client agreement, twenty representative document packages assessed by two auditors.

Gate At least 2.5 net hours per week, quality and cost confirmed

MONTHS 4 TO 6

Up to twenty users. The authorised audit RAG and the reporting templates go live, partner takes over support.

Gate Independent approval and monthly benefit tracking in place

MONTHS 7 TO 12

Model change and data export tested, contracts reviewed and growth loads simulated for forty users.

Gate Month twelve: cloud re-decided against on-premise on measured eligibility

If a gate fails, the manual workflow stays and the rollout stops. **A failed pilot costs EUR 12,000 and answers the question.**

Q8 IMPLEMENTATION STRATEGY

Every success measure has a number, an owner and a consequence

8

WHO GOVERNS IT, HOW IS SUCCESS MEASURED, AND WHAT IF PEOPLE RESIST?

AREA	WHO OWNS IT	HOW WE MEASURE SUCCESS	WHAT WE DO ABOUT FRICTION
Quality	Audit and quality lead	95 % correct evidence references, no critical decision error	Two senior auditors design the test set, so the bar is theirs
Security	External IT partner	No unauthorised result in negative-access tests	Access rights are cleaned before switch-on
Economics	Management	2.5 net hours per auditor per fully utilised week	Released time funds quality first, so nobody loses by reporting it
Adoption	Internal product owner	80 % of pilot users use it meaningfully each week	Weekly feedback decides what gets built next
Compliance	Privacy function	Signed ICT annex on file for every client in scope	Refusal is a normal answer; the manual method stays available

Management approves two things today: the pilot budget of EUR 12,000 and the name of the internal product owner. Everything else follows a gate.

The honest risk: experienced auditors may read this as a judgement on their speed. It is not, and the measures are built so that reporting a saving never costs anyone anything.

+ ADDITIONAL CONSIDERATION, RAG

Yes, build the audit RAG, because it makes every answer checkable

WHAT IS RAG, AND HOW DOES IT DIFFER FROM UPLOADING FILES TO CHATGPT?



1 APPROVED BASE	2 RETRIEVE	3 ANSWER	4 CITE	5 AUDITOR APPROVES
Only documents we may legitimately reuse	Permissions enforced before results are picked	Only the retrieved passages reach the model	Source, page and version shown with every claim	No evidence, no answer. Approval is recorded

The difference is not the algorithm. A chatbot upload also retrieves. Business RAG adds maintained sources, client separation, version control, deletion rules and a tested answer set.

SHOULD IT BE PART OF THE SOLUTION, AND HOW DO WE BUILD IT?

- Yes, from month four, once the pilot gates pass
- Index inside the EU, protected like the originals
- Model stays interchangeable after quality testing

WHAT DOES RAG NOT SOLVE?

- Retrieved passages still leave the internal network
- A correct citation is not a correct conclusion
- It improves traceability and does not protect confidentiality

EDHEC BUSINESS SCHOOL

OUR RECOMMENDATION

Deploy a controlled business cloud and prove it on a twelve-week pilot

EUR 77,424	EUR 119,232	EUR 41,808	EUR 12,000
Three-year cost	Capacity released	Net value	What we ask for today

APPROVE TODAY

The pilot budget and a named internal product owner.

REVIEW IN TWELVE WEEKS

Quality, access controls, eligible workload and measured time savings.

REOPENS THE DECISION

Eligible workload below 64 per cent, or a client requiring local processing.